

LLM-Based Document & Data Analyzer

1 What exactly is an LLM-based Document & Data Analyzer?

It is a system that:

- Reads **documents** (PDF, DOCX, TXT)
- Understands **tabular data** (CSV, Excel)
- Accepts **natural language questions**
- Returns:
 - Correct answers
 - Logical explanations
 - (Optionally) source references

 Important:

LLM **does NOT store knowledge** — it reasons on top of **your uploaded data**.

2 Problem this project solves (VERY IMPORTANT)

Real-world problem:

- Documents are **long**
- Data is **scattered**
- Humans:
 - Miss details
 - Make calculation mistakes
 - Waste time

Example:

200-page report + 5 CSV files

Question:

“What are the main risks and how did revenue change YoY?”

Manual → 

LLM-based system → 

3 What LLM actually does in this project

LLM is used for **THINKING**, not storing data.

Task	Done by
Understanding question	LLM
Breaking question into parts	LLM
Finding relevant text	Vector DB
Doing calculations	Python
Writing final answer	LLM

4 Why we cannot directly give documents to LLM

-  Token limit
-  Slow
-  Hallucination
-  Expensive

Solution:

We **index documents** instead of sending everything.

5 Core Concept: RAG (Retrieval-Augmented Generation)

Simple meaning:

“Search first → then think → then answer”

Flow:

User Question



Relevant Document Retrieval



LLM Reasoning on Retrieved Content



Final Answer

 But we will go **beyond basic RAG** later.

6 Types of Data in this Project

A. Unstructured Data (Text)

- PDFs
- DOCX
- TXT

Handled using:

- Text extraction
 - Chunking
 - Embeddings
 - Vector search
-

B. Semi-Structured Data (Tables)

- CSV
- Excel
- Tables inside PDFs

Handled using:

- Pandas
 - DuckDB / SQL
 - Python execution
-

7 Why we separate TEXT and TABLE data

Because:

- LLMs are **bad at math**
- Python is **perfect at math**

Data Type	Engine
-----------	--------

Text reasoning LLM

Numerical reasoning Python

Final explanation LLM

This makes system **accurate + reliable**.

8 What happens when user asks a question

Example question:

“Compare Q2 revenue growth and explain risks.”

Step-by-step:

1. LLM understands question
 2. Splits into:
 - Numerical part (revenue growth)
 - Textual part (risks)
 3. Retrieves:
 - Tables for numbers
 - Text for risks
 4. Python calculates numbers
 5. LLM explains results
-

9 What is Hallucination (and why we fight it)

Hallucination = LLM invents facts 

Why it happens:

- Missing data
- Weak prompts
- Overconfidence

Our prevention:

- Retrieval-only answers
 - Source verification
 - Validation step
-

10 Minimum Skills Needed (You already have most)

- ✓ Python basics
 - ✓ Pandas
 - ✓ API usage
 - ✓ Basic NLP understanding
- ✗ No deep ML training needed
 - ✗ No GPU needed
-

1 1 High-Level System Components

Component	Purpose
File Loader	Read documents
Chunker	Break text
Embedding Model	Convert text to vectors

Vector DB

Semantic search

Query Analyzer

Understand question

Reasoning Engine

Logic + math

Answer Generator

Final response

1 2 What makes THIS project interview-grade

- Multi-modal reasoning (text + data)
- Real-world relevance
- Anti-hallucination design
- Clean architecture
- Extensible system

1 3 What this project is NOT

- ✗ Simple chatbot
- ✗ ChatGPT wrapper
- ✗ One-file script

This is a **SYSTEM**, not a demo.

1 4 Project Outcome (End Goal)

By end, your system can answer:

- “Summarize this report”
- “Calculate average growth”
- “Compare two datasets”
- “Explain risks from documents”
- “Give source-backed insights”



STEP 2: SYSTEM ARCHITECTURE

(Document & Data Analyzer – Reasoning-Based LLM System)

1 High-Level Architecture (Bird's Eye View)

Think of the system as **6 clear layers**:

User

↓

UI Layer

↓

Query Understanding Layer

↓

Retrieval Layer

↓

Reasoning Layer

↓

Answer Generation Layer

2 UI Layer (Input & Output)

Responsibilities:

- Upload documents (PDF, CSV, Excel)
- Accept natural language questions
- Display answers + sources

Tools:

- Streamlit (best for start)
- OR FastAPI + frontend later

 UI does **no thinking**.

3 Document Ingestion Layer (Knowledge Creation)

What happens here:

- Files are loaded
- Text & tables are extracted
- Data is structured

For TEXT:

- PDF → text
- DOCX → text
- Chunked into small pieces

For TABLES:

- CSV / Excel → Pandas DataFrames

- Stored separately

📌 Important:

Ingestion happens **once**, not per question.

4 Query Understanding Layer (MOST IMPORTANT)

Purpose:

Understand **what kind of question** user is asking.

LLM decides:

- Is it:
 - Summary?
 - Numerical calculation?
 - Comparison?
 - Explanation?
- Which data is needed:
 - Text?
 - Tables?
 - Both?

Output

```
json
{
  "question_type": "comparison",
  "needs_text": true,
  "needs_table": true,
  "metrics": ["revenue", "growth"],
  "time_period": "Q2"
}
```

🔥 This is what makes project HARD.

5 Retrieval Layer (Finding Evidence)

A. Text Retrieval

- Convert question → embedding
- Vector search in DB
- Return top-k chunks

B. Table Retrieval

- Identify relevant DataFrame
- Filter rows/columns
- Send to reasoning layer

⚠️ LLM NEVER sees full data.

6 Reasoning Layer (Brain of System)

Split into 2 engines:

◆ Numerical Reasoning

- Python executes:
 - Averages
 - Growth rates
 - Comparisons

◆ Textual Reasoning

- LLM reads retrieved text
- Extracts insights
- Identifies causes, risks, explanations

7 Answer Validation Layer (ANTI-HALLUCINATION)

Before final output:

- Check:
 - Are numbers computed?
 - Are claims supported by retrieved text?
- If not:
 - Regenerate answer
 - OR say “Insufficient data”

🔥 This layer is rare & impressive.

8 Answer Generation Layer

LLM generates:

- Clear explanation
- Structured output
- Optional citations

Example:

Revenue grew by 12% in Q2, driven by increased sales in Region A.
Key risks include supply chain instability (Report.pdf, Page 18).

Project structure

```
doc_analyzer/
    ├── data/
    │   ├── raw_docs/
    │   └── processed/
    ├── ingestion/
    │   ├── load_text.py
    │   ├── load_tables.py
    │   └── chunking.py
    ├── embeddings/
    │   └── embed_store.py
    ├── query_engine/
    │   ├── query_parser.py
    │   └── intent_detector.py
    ├── reasoning/
    │   ├── text_reasoner.py
    │   ├── table_reasoner.py
    │   └── validator.py
    └── app.py
    └── requirements.txt
```

10 Data Flow (IMPORTANT FOR EXPLANATION)

1. Upload docs → ingestion
2. User asks question
3. Query parser classifies intent
4. Retriever fetches evidence
5. Reasoning engine processes
6. Validator checks correctness
7. Final answer shown

1 1 Why this architecture is STRONG

- ✓ Scalable
- ✓ Low hallucination
- ✓ Clear separation
- ✓ Extendable (images, audio later)